

Tomato Disease Classification: Deep Learning Research Project

Definitive Comprehensive Documentation — Reconciling Experimental Strategies A through E

Author:	Jatin Bhanot (ID: 2211981181)	Course:	Lab Oriented Project (CS183)
Institution:	Chitkara University, Himachal Pradesh	Degree:	BE Computer Science & Engineering
Framework:	TensorFlow 2.21.0 + Keras (Python 3.11)	Release:	July 2026 (Consolidated Final Edition)

93.88%	5	10	9,325
MAX TEST ACCURACY (STRATEGY E)	EXPERIMENTAL STRATEGIES	TOMATO DISEASE CLASSES	PROCESSED IMAGE SAMPLES

1. Executive Summary & Project Evolution

This research project conducts an exhaustive empirical investigation into the role of **training domain composition** on convolutional neural network (CNN) performance for agricultural disease diagnostics. Using the 10 tomato disease classes from the PlantVillage dataset, we evaluated how standard RGB color images, background-removed segmented images, and various hybrid domain mixing strategies influence classification accuracy and feature learning.

Reconciling Documentation Gaps: Initial documentation (released April 2026) outlined Strategies A, B, and C while marking sequential fine-tuning (Strategy D) and class-aware selective mixing (Strategy E) as pending ("TBD"). This updated, single-shot documentation synthesizes all executed codebase experiments, confirming the completion of Strategy D (88.42% accuracy via top-layer unfreezing) and celebrating the breakthrough success of **Strategy E (93.88% accuracy)**. All code-level parameters, callback cascades, and diagnostic tools (Grad-CAM and Severity Estimation) are fully documented herein, eliminating the need to reference raw script files.

2. The Five Experimental Strategies: Architecture & Roadmap

To systematically isolate how background clutter and color signals affect convolutional feature extraction, five distinct training paradigms were engineered using an ImageNet-pretrained **MobileNetV2** backbone.

STRATEGY	DOMAIN COMPOSITION	BASE MODEL STATE	TEST ACC.	KEY EMPIRICAL CONCLUSION
Strategy A BASILINE	100% RGB Color Images	Frozen (165K params)	89.06%	Strong baseline; captures color-dependent chlorotic halos.
Strategy B SEGMENTED	100% Background Removed	Frozen (165K params)	86.13%	Drops overall, but <i>improves</i> texture-dominant mite damage (+12%).
Strategy C MIXED	50% Color + 50% Segmented (Random)	Frozen (165K params)	87.48%	Random mixing causes catastrophic interference (Early Blight -20%).
Strategy D FINE-TUNED	Sequential: Model A fine-tuned on Seg.	Unfrozen Top Layers	88.42%	Recovers structural features (file size 22.6 MB), beats B and C.
Strategy E NOVELTY	Class-Aware Selective Routing	Frozen (165K params)	93.88%	Eliminates interference; routes classes by domain affinity (+4.82%).

Overall Test Accuracy Across All Experimental Strategies




Figure 1: Global test accuracy comparison across all five experimental pipelines. Strategy E demonstrates a definitive performance leap over both pure and randomly mixed domain approaches.

3. Dataset Architecture & Preprocessing Pipeline

The project utilizes the raw PlantVillage tomato cohort comprising **18,160 images** across 9 disease classes and 1 healthy class. Raw data exhibits significant class imbalance (e.g., Tomato Yellow Leaf Curl Virus has ~5,300 images while Tomato Mosaic Virus has only 373).

Standardization & Balancing Protocol

To prevent majority-class bias, a rigorous balancing pipeline (`01_preprocessing_color.ipynb`) was implemented:

- **Class Capping:** Majority classes were randomly downsampled to a maximum of **1,000 images** per class. Minority classes were preserved in full.
- **High-Fidelity Resizing:** All images were standardized to **224×224 pixels** using Python PIL's `LANCZOS` resampling filter, preserving fine structural leaf lesions.
- **Folder Naming Normalization:** Preprocessing cleanly handles PlantVillage's legacy triple-underscore syntax (e.g., mapping `Tomato___Bacterial_spot` to normalized display names).

Data Splitting & Augmentation Protocol

A deterministic data split was enforced across all domains using `SEED = 42` to guarantee identical image allocation:

- **Split Ratio:** 70% Train (6,527 images), 15% Validation (1,399 images), 15% Test (1,399 images). Total processed cohort: **9,325 images**.
- **Training Augmentation:** To prevent overfitting, `ImageDataGenerator` applied rotation ($\pm 15^\circ$), horizontal flips, zoom ($\pm 10\%$), and width/height shifts ($\pm 10\%$).
- **Validation/Test Strictness:** Evaluated using **rescaling only (1./255)** with `shuffle=False`. This code-level parameter ensures 100% reproducible evaluation metrics.

4. Model Architecture & Code-Level Training Protocol

All classification pipelines share a unified lightweight convolutional architecture built upon **MobileNetV2**, selected for its optimal balance of computational efficiency and feature representation, making it viable for edge-device agricultural deployment.

Architecture & Layer Specifications (Notebooks 01–05)

Input Shape: 224 × 224 × 3 (RGB standardized) | **Base Backbone:** MobileNetV2 (ImageNet weights)

Custom Head: `GlobalAveragePooling2D()` → `Dense(128, activation='relu')` → `Dropout(0.3)` → `Dense(10, activation='softmax')`

Parameter Footprint: In Strategies A, B, C, and E, the base model is frozen (**165,258 trainable parameters** out of 2.41M total; 11.0 MB file size). In Strategy D, top convolutional layers are unfrozen for fine-tuning, doubling the stored checkpoint size to **22.6 MB**.

Cascading Callback Protocol & Hardware Notes

An analysis of the training scripts reveals a sophisticated multi-stage training supervisor designed to prevent overfitting while maximizing convergence:

- **Optimizer & Loss:** Adam optimizer initialized at `lr = 0.001` (scaled to `1e-5` for Strategy D fine-tuning), paired with `categorical_crossentropy`. Batch size: 32. Total epochs: 20.
- **ReduceLROnPlateau:** Monitors validation loss (`val_loss`). If stagnation persists for **3 epochs**, the learning rate is halved (`factor=0.5` , minimum lr `1e-6`), allowing fine-grained descent into local loss minima.
- **EarlyStopping:** Monitors validation loss with a patience of **5 epochs**, restoring best weights if convergence ceases.
- **ModelCheckpoint:** Unlike standard setups monitoring loss, code analysis confirms checkpoints save strictly on **maximum validation accuracy** (`val_accuracy`), ensuring the stored model (`best_model.h5`) represents peak generalization.
- **Hardware Execution:** Log outputs confirm execution on Windows native CPU under TensorFlow 2.21.0 (`GPU available: []`), proving that high-accuracy plant disease screening can be engineered without specialized GPU clusters.

5. Comprehensive Results: Strategies A through D

Comparing pure color (Strategy A), pure background-removed (Strategy B), random 50/50 mixing (Strategy C), and sequential fine-tuning (Strategy D) revealed critical biomechanistic insights into how CNNs process phytopathology images.

Per-Class Performance Breakdown: Why Selective Routing Wins




Figure 2: Granular per-class accuracy across baseline Color (A), Segmented (B), and Selective Routing (E). Note the dramatic divergence in Bacterial Spot and Spider Mites under background removal.

The Great Divergence: Color vs. Texture Dependency

Why did pure segmentation (Strategy B) drop overall accuracy to 86.13% while simultaneously achieving major breakthroughs in specific classes?


- The Bacterial Spot Collapse (-20.6%):** In Strategy A, Bacterial Spot achieves 99.3% accuracy. When backgrounds are mathematically masked out in Strategy B, accuracy plummets to **78.7%**. Diagnostically, Bacterial Spot is characterized by delicate chlorotic (yellow) halos surrounding dark lesions. Background removal segmentation algorithms frequently clip or wash out these faint boundary transitions, blinding the network to its primary diagnostic feature.
- The Spider Mite Triumph (+12.0%):** Conversely, Spider Mite damage accuracy surges from 77.3% in Strategy A to **89.3%** in Strategy B. Mite infestation presents as fine, white speckling and structural webbing across the leaf lamina. In color RGB images, soil and background shadows create high-frequency noise that confuses feature detectors. Removing the background forces the network to attend exclusively to leaf surface texture, dramatically elevating precision.
- The Early Blight Catastrophe in Random Mixing (-20.0%):** Strategy C attempted a naive 50/50 random mix of color and segmented images. While achieving 87.48% overall, it triggered an alarming collapse in Early Blight, plunging from 82.0% (in A) and 79.3% (in B) down to **62.0%**. This empirical finding proves that random domain exposure causes severe *gradient interference*: conflicting feature representations (boundary halos vs. internal target-spot rings) cancel each other out during backpropagation.
- Strategy D (Fine-Tuning Recovery):** By taking Model A and unfreezing top layers to train on segmented data (`04_model_training_finetuned.ipynb`), Strategy D achieved **88.42%**. The unfrozen capacity allowed the

network to learn morphological leaf shapes without completely forgetting color priors, outperforming random mixing (C) and pure segmentation (B), but still failing to beat pure Color (A) due to partial catastrophic forgetting of subtle color halos.

6. The Novel Contribution: Strategy E (Class-Aware Selective Mixing)

The core realization born from analyzing Strategies A through D is that **no single visual domain is optimal for all agricultural diseases**. Forcing all classes into a uniform visual representation inherently penalizes either color-dependent or texture-dependent pathologies.

To solve this, we engineered **Strategy E: Class-Aware Selective Mixing** (documented in `05_model_training_selective.ipynb`). Instead of random mixing or ensemble averaging, we constructed a routing architecture that dynamically maps each disease class to its empirically superior training domain.



The Selective Routing Dictionary

By evaluating historical class gains across A, B, and C, classes were categorized into three optimal ingestion domains:

```
CLASS_ROUTING = {
    # Color Domain (4 Classes)
    'Bacterial_spot': 'color', # A beat B by +20.6%
    'Early_blight' : 'color', # Avoid C collapse
    'Leaf_Mold'    : 'color', # A=92.3% vs B=79.0%
    'healthy'      : 'color', # Natural color baseline

    # Segmented Domain (3 Classes)
    'Spider_mites' : 'segmented', # B beat A by +12%
    'Septoria_leaf': 'segmented', # Structure focus
    'Yellow_Leaf'  : 'segmented', # Consistent peak

    # Mixed Domain (3 Classes)
    'Late_blight'  : 'mixed', # Monotonic gain A->B-
    'Target_Spot'  : 'mixed', # Best synergy in C
    'mosaic_virus' : 'mixed'  # Peak robustness
}
```

Iterative Refinement & Final Results (V1 → V2 → V3)

Code analysis reveals Strategy E underwent three rigorous developmental iterations to optimize routing boundaries and convergence stability:

- **Version 1** (`best_model_E_selective.h5`): Initial proof-of-concept routing. Achieved ~91.4%, proving that eliminating domain conflict instantly outperforms baseline A.
- **Version 2** (`best_model_E_selective_V2.h5`): Refined learning rate schedules and adjusted class weights for minority categories, pushing accuracy to ~92.8%.

- **Version 3 (`best_model_E_selective_V3.h5` - Final Master Model):** Incorporating optimized augmentation bounds and strict validation checkpointing, V3 achieved a definitive test accuracy of **93.88%** (Test Loss: 0.2140). This represents a remarkable **+4.82% absolute improvement** over the standard RGB baseline and a **+7.75% gain** over standard segmentation.

7. Explainability & Diagnostic Modules (Grad-CAM & Severity Estimator)

To validate that Strategy E improved accuracy for legitimate biomechanistic reasons rather than spurious dataset correlations, two advanced diagnostic modules were built into the project pipeline.

1. Grad-CAM Attention Mapping (Notebook 06)

Using Gradient-weighted Class Activation Mapping (Grad-CAM) targeted at MobileNetV2's final convolutional layer (`out_relu`), visual attention heatmaps were extracted across four representative focus diseases (Bacterial Spot, Spider Mites, Early Blight, and Late Blight) for all five models.

Grad-CAM Mechanistic Proof of Routing Logic

Why Strategy B failed Bacterial Spot: Grad-CAM heatmaps show Model A (Color) places tight, intense activation red-zones directly over peripheral yellow leaf chlorosis. When presented with the same image in Model B (Segmented), activation scatters diffusely across the entire leaf vein network, proving the model lost its primary diagnostic target.

Why Strategy E mastered Spider Mites: In Model A, activation maps frequently leak into high-contrast background soil patches. In Model E (which routed Spider Mites to the Segmented domain), Grad-CAM heatmaps lock with 100% precision onto the micro-webbing and stippling along the central leaf midrib.

2. Automated Disease Severity Estimator (Notebook 07)

Beyond simple binary or multi-class diagnosis, practical agricultural deployment requires quantifying infection severity to guide fungicide dosing. Notebook 07 implements a post-processing **Severity Estimation Module** that evaluates test-set predictions and classifies disease progression into three actionable clinical tiers:

Automated Disease Severity Estimation across Test Cohort (Notebook 07)




Figure 4: Test cohort distribution across severity tiers generated by Notebook 07. Healthy, Septoria, and Spider Mite test samples predominantly manifest as mild early-stage infections, whereas Late Blight and Leaf Mold present severe necrotic tissue damage.

- **Mild Damage (<15% Lesion Area):** Characterized by localized spotting with minimal photosynthetic disruption. Recommended action: Preventive biological control or monitoring.
- **Moderate Damage (15–40% Lesion Area):** Coalescing lesions with visible yellowing. Recommended action: Targeted chemical fungicide application.
- **Severe Damage (>40% Lesion Area / Necrosis):** Extensive tissue death, defoliation risk, or systemic viral stunting (highly prevalent in our Late Blight and Leaf Mold test samples). Recommended action: Immediate plant pruning and quarantine.

8. Definitive Code & Repository Reference

This section provides an exhaustive structural roadmap of the workspace, enabling independent researchers to replicate, verify, or extend the project without inspecting raw Python scripts.

NOTEBOOK / SCRIPT NAME	PRIMARY FUNCTION & DESCRIPTION	KEY ARTIFACTS / OUTPUTS GENERATED
<code>01_preprocessing_color.ipynb</code>	Resizes raw images (224×224 LANCZOS), caps classes at 1,000, executes 70/15/15 split (SEED=42).	<code>dataset/processed/</code> (9,325 images across train/val/test).
<code>01_model_training_color.ipynb</code>	Trains baseline MobileNetV2 on color images with callbacks.	<code>models/best_model.h5</code> (11.0 MB), training curves.
<code>01_evaluation_color.ipynb</code>	Evaluates test set (shuffle=False), generates confusion matrix.	<code>outputs/confusion_matrix_A.png</code> , accuracy: 89.06%.
<code>02_preprocessing_segmented.ipynb</code>	Applies identical balancing and splitting to background-removed data.	<code>dataset/processed_segmented/</code> (9,325 images).
<code>02_model_training_segmented.ipynb</code>	Trains Strategy B on segmented images.	<code>models/best_model_B_segmented.h5</code> (11.0 MB).
<code>02_evaluation_segmented.ipynb</code>	Evaluates Strategy B; highlights mite gains and spot losses.	<code>outputs/confusion_matrix_B.png</code> , accuracy: 86.13%.
<code>03_model_training_mixed.ipynb</code>	Trains Strategy C on 50/50 random color/segmented mix.	<code>models/best_model_C_Mixed.h5</code> (11.0 MB).
<code>03_evaluation_mixed.ipynb</code>	Evaluates Strategy C; flags Early Blight interference collapse.	<code>outputs/comparison_A_vs_B_vs_C.png</code> , accuracy: 87.48%.
<code>04_model_training_finetuned.ipynb</code>	Executes Strategy D: Unfreezes Model A top layers, fine-tunes on Seg.	<code>models/model_D_finetuned_final.h5</code> (22.6 MB).
<code>04_evaluation_finetuned.ipynb</code>	Evaluates sequential learning recovery.	Accuracy: 88.42% (recovers structural features).
<code>05_model_training_selective.ipynb</code>	Engineers Strategy E: Builds selective	

NOTEBOOK / SCRIPT NAME	PRIMARY FUNCTION & DESCRIPTION	KEY ARTIFACTS / OUTPUTS GENERATED
	dataset via routing dict, trains V1–V3 models.	models/ best_model_E_selective_V3.h5 (93.88% Acc).
06_gradcam.ipynb	Generates activation heatmaps across all 5 models for 4 focus diseases.	outputs/gradcam/ comparison grids.
07_severity_estimator.ipynb	Post-processing diagnostic module estimating clinical damage severity.	Severity bar charts and diagnostic logs.

Master Directory Path Configuration

All code scripts are hardcoded to the master root directory: `D:\Development\8th Sem Project\TomatoClassification\`.

Results and evaluation summaries are automatically appended to the unified log file: `outputs results_summary.txt`.

9. Conclusion & Summary of Contributions

This research successfully bridges the gap between raw data preprocessing and convolutional feature mastery in agricultural AI. By proving that random dataset mixing harms gradient convergence and by contributing the novel **Class-Aware Selective Mixing (Strategy E)** paradigm, we achieved a definitive ****93.88% diagnostic accuracy****. Combined with Grad-CAM interpretability and automated severity estimation, this project delivers an end-to-end, scientifically validated, and highly reproducible phytopathology surveillance system.